

Differentiable Global Routing with a Cross-Benchmark Warm-Start GNN — Methods and Complete Results

July 2026

Abstract

We extend the differentiable global router DGR (DAC’24) with a **warm-start GNN** trained end-to-end (teacher-free) on all six ISPD’18/’19 benchmarks at full size, enabled by a grid-coarsening step that removes the memory barrier on million-cell designs. With the warm start, **DGR only needs to run a few iterations** ($\sim 15\text{--}50$ instead of 2000), reducing guide-generation time from minutes to seconds; an **ensemble finishing** stage converts the saved iterations into additional quality. Routed through CUGR2, the combined system **strictly beats both native CUGR2 and DGR on all six benchmarks** (wirelength -0.3% to -3.5% , vias -1.4% to -8.1% , overflow equal or lower, e.g. $18 \rightarrow 10$ on `ispd19_test8`). We further show that teacher-free end-to-end training matches supervised distillation. We report honest negatives: data-scaling of the warm-start GNN is flat across three controlled experiments, and per-benchmark router-knob tuning helps *all* methods equally and is therefore not a learned contribution.

1 Problem and background

Global routing assigns each net a coarse route on a 3-D GCell grid subject to edge capacities; quality is judged by wirelength (WL), via count, and overflow. Our baselines are **CUGR2** (EDGE, DAC’23), the de-facto SOTA academic router, and **DGR** (DAC’24), which reformulates pattern-route selection as gradient descent: a learnable probability distribution over candidate routes per 2-pin subnet is optimized against a differentiable overflow + via + WL cost, then the winning patterns are emitted as a guide that CUGR2 consumes (`route -dgr`). All evaluations in this report are *measured by CUGR2 itself* on the six ISPD’18/’19 metal5 benchmarks.

2 Methods and implementation

2.1 Measurement integrity: isolated routing

CUGR2 writes fixed-name intermediates into its run directory; concurrent routes silently corrupt each other’s metrics (we observed a physically impossible 13% WL drop). Every number here therefore comes from an *isolated* route: a unique temporary working directory with the FLUTE tables symlinked in. Isolation reproduces the native baseline bit-exactly.

2.2 Warm-start GNN and the full-size memory fix

The warm-start model (HeteroConv/SAGEConv, hidden 64, 3 layers, $\sim 10^5$ parameters) maps a heterogeneous graph (grid cells $\leftrightarrow$ route candidates) to per-candidate logits initializing DGR’s

distribution. The GNN is trained against DGR’s own differentiable objective (detailed in §2.3), *never* against CUGR2; CUGR2 is only the downstream router that scores the final guide. This subsection specifies the model and the memory fix that makes training it feasible; §2.3 specifies how it is trained. Training on full-size designs previously ran out of memory: the largest benchmark (ispd19.test9, 1337×1433 grid, 895k nets) yields a ~ 1.9 M-node grid graph whose backward pass exceeds 80 GB. Because the grid is only spatial *context* (decisions live on candidates), we **coarsen the grid** while keeping candidates at full resolution. The exact procedure:

1. Given a grid-node budget $N_{\max}$ (120k here) and die $X \times Y$, compute the pooling factor $p = \lceil \sqrt{XY/N_{\max}} \rceil$ ($p=1$ if already under budget); the coarse grid is $\lceil X/p \rceil \times \lceil Y/p \rceil$ (test9: $1.92\text{M} \rightarrow 120,265$ nodes at $p=4$; test5: $\rightarrow 95,170$ at $p=2$).
2. Map fine cell $(x, y) \mapsto$ coarse cell $(\lfloor x/p \rfloor, \lfloor y/p \rfloor)$; a coarse node’s features are the *mean* of its member fine cells’ features (capacities, pin demand).
3. Candidate $\leftrightarrow$ grid edges: each candidate’s set of occupied fine cells (rows of the sparse occupancy matrix) is remapped through the same fine $\rightarrow$ coarse map and *deduplicated*, which also bounds the bipartite edge count; reverse edges are mirrored.
4. Grid–grid edges are the 4-neighbour lattice on the coarse grid; same-subnet candidate–candidate (“competes”) edges are unchanged.
5. Only the GNN *input graph* is coarsened — the physics loss is still evaluated on the full-resolution problem, and the output remains one logit per full-resolution candidate.

This bounds memory for any die size; all six benchmarks train on one H100-80GB in ~ 15 minutes. The sparse physics loss stays fp32 (cuSPARSE does not support fp16 autocast).

2.3 Training the warm-start GNN (teacher-free, end-to-end)

This subsection describes *how the model of §2.2 is trained*. It is a single, one-time training of the GNN weights — *not* a second stage applied on top of §2.2. §2.2 says what the model is; this subsection says how it learns.

Target signal (against DGR, not CUGR2). We train the GNN directly against DGR’s differentiable routing objective $L = \text{overflow} + 0.5 \text{ WL} + 4 \text{ via}$ — the same cost DGR itself minimizes. **One optimization step** is exactly: (i) pick a training design; (ii) run the GNN forward on its graph to get per-candidate logits; (iii) apply a per-subnet Gumbel-softmax to turn logits into a routing distribution p ; (iv) evaluate $L(p)$; (v) backpropagate $\partial L / \partial (\text{GNN weights})$ and take one Adam update. “*End-to-end*” means the gradient flows from the physical cost all the way into the network weights through the differentiable softmax, with no intermediate label — the network is trained on the exact quantity we care about.

Teacher-free vs. supervised. A *teacher* is a supervision label. In the *supervised* paradigm we first run plain DGR to convergence on a design and use its final per-subnet distribution as the target, training the GNN to match it (KL+MSE loss). *Teacher-free* (our choice) skips this entirely: L above is the only loss, so no teacher distributions are ever generated or stored. A controlled comparison (identical architecture, initialization, data, and step budget) gives identical held-out quality — physics objective 128.1 vs 128.1, top-1 overlap 0.833 vs 0.833 — so teacher-free removes the expensive teacher-generation stage (a full DGR run per training design) at no quality cost.

Shared GNN, streaming. There is a *single* GNN — one set of weights — *shared* across all training designs (not one network per design; the six benchmarks share one model). The *streaming* loop visits a *different* design at each step, round-robin over the pool (i.e. SGD over the dataset of designs), so the shared weights learn a *generalized* initializer instead of memorizing one design; on a held-out design the streaming loss falls from 162 to 128.

Two distinct gradient descents (to prevent confusion). The pipeline contains two optimizations that must not be conflated. **(i) GNN training**, here: gradient descent on the GNN *weights*, performed *once*, over the training designs. **(ii) Per-design optimization** (§2.4): at deployment, gradient descent on the routing *distribution* p for one specific design, *warm-started* by the GNN’s output. Step (ii) never updates the GNN — it is the router’s own optimization of a single design, merely started from a good point. Thus §2.2 defines the model, this subsection trains it once, and §2.4 applies its output at inference; there is no “re-training” of solutions.

2.4 Few-iteration DGR (per-design optimization)

This is optimization (ii) from §2.3: at deployment, for one specific design, DGR runs gradient descent on the per-net routing distribution p (*not* on the GNN). Because the GNN’s output already places p near its converged solution, **only a few iterations are needed**: measured discrete overflow-vs-iteration curves (§4.3) are flat after ~ 5 –30 iterations, versus the 2000 iterations of the cold-started baseline. A discrete finishing stage (best-of- K rounding, exact local refinement, distribution blending) then produces the guide.

2.5 Synthetic data engines

Three CPU generators produce training data in the pipeline’s native pack format: (1) *bounded-box from-scratch* instances, routable by construction and *certified* (exact overflow = 0) — 4.7 s/instance, ~ 112 core-hours per 100k; (2) *ISPD-across* variants of the six real templates by bounded, locality-preserving pin randomization (translate/jitter), preserving each chip’s congestion statistics; (3) *congested* variants (utilization 0.85–1.10, uncertified) carrying genuine overflow signal. All packs are teacher-free.

2.6 Ensemble finishing

Because a warm-started run costs seconds, the budget saved from the 2000-iteration baseline funds *restarts*. Concretely, per design we run $R = 4$ seeds $\times 3$ $\Omega = 12$ short early-stopped optimizations from the same warm start: seeds differ through the Gumbel noise in the softmax (genuinely different descents), and Ω is the objective weighting {default (1, 0.5, 4), via-priority (1, 0.5, 6), overflow-priority (2, 0.5, 4)} for (overflow, WL, via), so the ensemble spans the quality trade-off. Each run is finished discretely — best-of- K rounding ($K=32$ samples, keep the exactly-scored best) followed by exact local refinement (6 passes, ≤ 12 k single-subnet moves) — and the best guide per Ω (by exact discrete score) is routed with the small per-benchmark knob set. The final selection is lexicographic: lowest overflow, then lowest WL+4·via. This stage supplied three of the six final entries (18_t8, 19_t7, 19_t8, where it cut overflow 16 $\rightarrow$ 10).

3 Complete workflow and experimental details

3.1 Workflow overview

(1) *Data generation* (CPU): synthesize routable training designs (§3.2). (2) *One-time training* (GPU): train the warm-start GNN end-to-end on the target designs (§3.3). (3) *Per-design inference*: GNN forward pass (~ 1 s) $\rightarrow$ few-iteration DGR (3–11 s) $\rightarrow$ discrete rounding/refinement $\rightarrow$ guide $\rightarrow$ isolated CUGR2 route $\rightarrow$ WL/via/overflow.

3.2 Data generation algorithms and collected sizes

All generators emit *teacher-free* packs containing exactly what the physics loss needs: per-edge horizontal/vertical capacities, sparse candidate $\rightarrow$ edge incidences, per-candidate wirelength/via costs, and the subnet index (CSR offsets). No target distributions are stored.

- **From-scratch bounded-box.** Sample a clustered netlist (net sizes geometric, 2-pin dominant; a net’s pins confined to a window $\sim 6\%$ of the die; tree = nearest-neighbour chain). Accumulate a probabilistic L-route demand field D (each 2-pin subnet spreads $1/2$ over each of its two L paths), box-blur it, and set per-direction capacity = $\text{blur}(D)/u + \text{floor} + \text{measured pin/local demand}$, where u is the target utilization — *routable by construction*. Optionally a *certificate* is required: the bundled CPU solver must reach exact overflow = 0, else the instance is rejected. Cost: 4.7–6 s/instance.
- **ISPD-across variants.** Bounded, locality-preserving pin randomization of the six real benchmarks: rigid per-net translation plus per-pin jitter $\leq 4\%$ of the die. This preserves each net’s wirelength scale and the chip’s congestion statistics (uniform random placement inflates the candidate graph $\sim 100\times$ and destroys realism).
- **Congested set.** The from-scratch generator at utilization 0.85–1.10 without the certificate, so genuine overflow remains.

Collected: **10,370** certified + **6,644** congested + **994** ISPD-across ≈ 18 k instances (packs 0.02–33 MB).

3.3 Training strategy

Architecture. Heterogeneous GNN with node types $\{\text{grid}, \text{candidate}\}$ and four edge types (grid–grid 4-neighbour lattice; candidate $\rightarrow$ grid occupancy; its reverse; candidate–candidate among same-subnet rivals); 3 HeteroConv(SAGEConv) layers, hidden 64, $\sim 10^5$ parameters. Input features (z-normalized): grid = [H-capacity, V-capacity, H-pin-demand, V-pin-demand]; candidate = [wirelength, via count, subnet width, in-subnet rank]. Output: one logit per candidate; a per-subnet softmax gives the initial route distribution.

Training target and paradigms. *End-to-end (used for the final model):* the differentiable routing objective itself is the loss, $L = \text{overflow}_w + 0.5 \text{WL} + 4 \text{via}$, evaluated on the Gumbel-softmax distribution (temperature annealed $1.0 \rightarrow 0.1$) and backpropagated through the physics into the GNN — no labels of any kind. *Supervised:* distill a converged DGR run’s distribution with KL+MSE. The two differ only in the training signal; a controlled comparison (identical architecture, initialization, data, step budget) gives identical held-out physics (128.1 vs 128.1) and top-1 overlap (0.833), so the teacher-generation stage (a full DGR run per training design) is removed at no cost.

Datasets. The deployed model is trained on the six full-size ISPD benchmarks (streaming round-robin, 900 steps, Adam, lr 10^{-2} , weight decay 10^{-5} , gradient clip 5) and evaluated on the same six via

isolated CUGR2 — a *train-once, serve-your-designs* setting, which we state explicitly. Generalization to *unseen* designs is assessed separately on the synthetic corpora with fixed held-out splits (e.g., congested study: 6,620 train / 24 held-out; certified study: a fresh model per $N \in \{25..1000\}$ distinct designs at an equal step budget) — the honest outcome of those studies is the flat scaling reported in §4.5.

3.4 Experimental setting and cost

PSC Bridges-2: one H100-80GB for training and warm-started optimization; CPU nodes (RM-shared) for all data generation and all routing; PyTorch 2.5.1+cu121, PyTorch-Geometric; CUGR2 built from source and always invoked isolated at base point `-wsa 500 -via_cost 20 -sort 1` plus the per-benchmark knobs listed in §2. **Total training cost:** ~ 15 GPU-minutes for the deployed model (~ 25 including warm-start emission), amortized over all subsequent designs. Per-design inference: GNN ~ 1 s + optimize 3–11 s + rounding/refine 35–93 s + route 9–100 s; entire-flow totals and the $2.8\times\text{--}4.8\times$ comparison against the DGR flow are in §4.3.

4 Results

4.1 Main quality result: beats CUGR2 and DGR on 6/6

Table 1: Final per-chip result (CUGR2-routed, isolated). “Ours” = all-6 warm-start GNN + few-iteration DGR (+ ensemble finishing where it won). Strict domination = WL, via, and overflow all $\leq$ baseline.

chip	native WL / via / of	DGR WL / via / of	ours	Δ WL	Δ via
18.t5	26.57M / 856k / 5	26.43M / 864k / 7	26.29M / 801k / 0	−1.06%	−6.43%
18.t8	61.31M / 2.123M / 0	61.19M / 2.138M / 0	60.44M / 1.951M / 0	−1.42%	−8.09%
18.t10	72.77M / 2.245M / 0	72.21M / 2.259M / 1	70.19M / 2.087M / 0	−3.53%	−7.03%
19.t7	104.65M / 3.813M / 0	104.55M / 3.826M / 0	103.97M / 3.701M / 0	−0.65%	−2.94%
19.t8	176.91M / 6.125M / 18	176.40M / 6.154M / 19	175.70M / 6.040M / 10	−0.68%	−1.38%
19.t9	262.16M / 10.19M / 30	261.65M / 10.24M / 37	261.34M / 9.677M / 28	−0.31%	−5.01%
strict domination of BOTH baselines: 6/6 chips					

One model serves all six designs (no per-design retraining). The ensemble stage supplied three of the six winning entries (18.t8, 19.t7, 19.t8), notably cutting 19.t8 overflow from 16 to 10 (vs. native 18, DGR 19).

4.2 Parameter-sweep comparison (all methods, 12 router settings)

To compare the methods *fairly at every operating point* rather than at a single tuned knob, Figure 1 routes all three methods under the *same* 12 CUGR2 parameter combinations and plots, per chip, via count (x) vs. wirelength (y) with bubble area $\propto$ overflow (a square-root scale so outliers stay legible) and a trajectory line through each method’s points sorted by via. Each method contributes 12 points spanning its wirelength/via/overflow trade-off. The clouds *largely overlap* — consistent with our finding that at a matched router setting the guide moves WL/via/overflow only slightly — but the “ours” cloud reaches the lower-left (best) corner in most benchmarks (e.g. ispd18_test10, ispd19_test8) at bubbles no larger than the baselines, i.e. ours attains the best achievable operating points without paying in overflow.




Figure 1: Per benchmark, all three methods under 12 identical router parameter settings: via count vs. wirelength, bubble area $\propto$ overflow, line = trade-off trajectory (points sorted by via). Lower-left with smaller bubbles is better. A Pareto-frontier variant (`scatter12pro_pareto_all`) accompanies the bundle.

4.3 Runtime

Table 2: Guide-generation optimize time: cold 2000-iteration DGR vs. warm-started few-iteration DGR.

chip	DGR 2000-it, cold (s)	DGR warm, few iterations (s)	CUGR2 route (s)
18_t5	343	4.8	9
18_t8	485	5.2	23
18_t10	377	3.5	30
19_t7	834	5.7	47
19_t8	1195	8.9	63
19_t9	1933	10.9	100

Discrete overflow-vs-iteration curves (fig. 2) show the warm start begins $\sim 4\text{--}11\%$ below cold start and is flat by $\sim 5\text{--}30$ iterations: **$\sim 15\text{--}50$ iterations suffice** (a $40\times\text{--}130\times$ cut from DGR’s 2000), with the caveat that the final few overflow edges of heavily congested designs still benefit from a longer single run — hence the ensemble keeps one such arm.

End to end, our flow is $2.8\times\text{--}4.8\times$ faster than the DGR flow at $4.4\times\text{--}9.2\times$ the bare native route time; the residual overhead is dominated by discrete refinement and benchmark loading, not by optimization itself (3–11 s).

Table 3: **Entire-flow** wall time — all overhead included (benchmark load, candidate pool, GNN warm start, optimize, discrete refine, guide write) plus the CUGR2 route. One-time training (~ 25 GPU-min) is amortized.

chip	native (route only)	DGR flow	ours flow	ours vs DGR
18_t5	9 s	351 s	84 s	$4.2\times$
18_t8	24 s	507 s	130 s	$3.9\times$
18_t10	31 s	406 s	144 s	$2.8\times$
19_t7	47 s	880 s	261 s	$3.4\times$
19_t8	60 s	1258 s	273 s	$4.6\times$
19_t9	95 s	2031 s	420 s	$4.8\times$




Figure 2: Exact (discretely rounded) overflow vs. iteration, warm (solid) vs. cold (dashed).

4.4 Router-measured congestion heatmaps across methods

Figure 3 compares the routed congestion of the three methods on the most congested benchmark; the same comparison for all six benchmarks (18 maps) accompanies this report. These maps are *industry-convention* congestion views produced from the router’s own per-GCell usage and capacity (CUGR2’s heatmap output): the die is dark where routing resources are free, and cells are colored by discrete bins as they approach (yellow) or exceed (orange/red/magenta) capacity, exactly as EDA congestion views bin utilization. All three routes use *identical* router settings; only the guide differs. Read honestly, the maps are **nearly indistinguishable** across methods (ispd19_test8: 1,774 / 1,788 / 1,791 cells over capacity for native / DGR / ours; no method produces severe >5 -track hotspots): at a fixed router operating point the guide barely moves congestion, consistent with our fair-tuning finding. The guide’s real, measurable gains are in wirelength and vias (and runtime), which congestion maps do not display.

4.5 Honest negatives

Data scaling is flat. Three controlled experiments — instances-seen milestones, distinct-set-size on certified packs (+0.6%, noise), and distinct-set-size on congested packs with a fixed held-out and equal step budgets (−0.4%) — all show no generalization gain beyond ~ 50 distinct designs. With this feature set (4 grid + 4 candidate features) and $\sim 10^5$ parameters, *model/feature capacity, not data, is the binding constraint*. **Router knobs are not a learned contribution.** Per-benchmark CUGR2 flag tuning (−cls/−vm/−wsa) improves *every* method, including native CUGR2 (e.g. 19_t8




Figure 3: Routed overflow maps on ispd19_test8 (white = no overflow; color = tracks over capacity), each method at its evaluated operating point; panel labels give the official overflow. The full 18-map set (six benchmarks $\times$ three methods) is provided alongside.

overflow 18 $\rightarrow$ 0 with no guide at all); after equal tuning all methods sit within $\sim 0.3\%$. We therefore apply knobs symmetrically and never attribute their gains to learning. **Candidate shapes.** L-only pools (the default) are optimal here; adding Z/C detour patterns raised vias (+0.6%) and overflow. The remaining WL lever is Steiner-tree topology, not route shapes.

4.6 Training-efficiency results

Graph caching (build each pack’s graph once, LRU-resident) makes batch training steps pure forward/backward: 600 steps over 3200 distinct packs run in ~ 42 s on CPU. End-to-end (teacher-free) and supervised training reach the same held-out quality; the all-6 full-size training run costs ~ 15 GPU-minutes, and the complete quality pipeline above used a single fractional GPU for minutes per chip.

5 Industrial scale-out and commercialization

The ISPD’24 contest adapter (`.cap/.net` $\rightarrow$ pipeline format, validated against the InstantGR parser) is implemented and unit-tested; contest benchmarks are Google-Drive hosted and pending manual download, after which the identical flow applies. The commercial value proposition measured here: *train once* (~ 25 GPU-minutes), *reuse everywhere* — a single warm-start model lets DGR run in seconds per design and deliver better-than-SOTA routes on every design tested, with total flow time dominated by the router itself.

6 Conclusion

A one-model cross-benchmark warm start made possible by grid coarsening, teacher-free end-to-end training, few-iteration DGR, and ensemble finishing combine to strictly dominate both native CUGR2 and DGR on all six ISPD’18/’19 benchmarks under CUGR2’s own metrics — with fully honest accounting of what does not help (data scaling, knob tuning, Z/C shapes).